

# CSE344: Computer Vision

## Assignment-1

Divyajeet Singh (2021529)

February 21, 2024

### Theory

#### Question 1.

- (a) The given problem of classifying papaya images into the classes ‘sweet’ and ‘not sweet’ is a binary classification problem. The Mean Squared Error function is non-convex for binary classification. Thus if a binary classifier is trained with MSE Loss, it is not guaranteed to converge to the minimum of the loss. Secondly, using MSE assumes an underlying Gaussian distribution<sup>1</sup>, which is not valid in case of binary classification, since it is modeled with a Bernoulli distribution.

- (b) For a single training example, the Binary Cross Entropy loss function is given by (assuming  $\hat{y}$  is a valid probability in  $[0, 1]$ )

$$\mathcal{L}(\hat{y}) = -y \log_2(\hat{y}) - (1 - y) \log_2(1 - \hat{y})$$

- (c) On a negative example (where  $y = 0$ ), if the predicted value is  $\hat{y} = 0.9$ , the value of loss is given by

$$\mathcal{L}(0.9) = -0 \log_2(0.9) - (1 - 0) \log_2(1 - 0.9) = -\log_2(0.1) \approx 3.322$$

- (d) Now, we calculate the mean loss for 3 examples for the given ground truth and predicted values

$$\begin{aligned}\mathcal{L}(\hat{Y}) &= -\frac{1}{3} \sum_{i=1}^3 y_i \log_2 \hat{y}_i + (1 - y_i) \log_2(1 - \hat{y}_i) \\ &= -\frac{1}{3} (\log_2(0.1) + \log_2(0.8) + \log_2(0.3)) \\ &\approx \frac{1}{3} (3.322 + 0.322 + 1.737) = \frac{5.381}{3} \approx 1.793\end{aligned}$$

- (e) For a learning rate  $\alpha$  and a trainable weight  $W$ , the update formula while using  $L_2$ -regularization (hyperparameter  $\lambda$ ) is given by

$$W_{\text{new}} \leftarrow W_{\text{old}} - \alpha \left( \frac{\partial L_{BCE}}{\partial W} + 2\lambda W_{\text{old}} \right)$$

The  $L_2$ -regularization applied in model  $A$  penalizes very large weights by adding some proportion of the weights themselves to the gradient update. Hence, we expect model  $A$ 's weights to be relatively smaller than model  $B$ 's final trained weights. However, this technique (usually) reduces overfitting, as model  $A$  will be encouraged to learn simpler patterns than model  $B$ .

- (f) KL-divergence is a measure of ‘distance’ or difference between two probability distributions. KL-divergence measures the relative entropy of two distributions. On the other hand, cross entropy can be thought of as the total entropy between them. The formula for KL-divergence and cross-entropy are as follows

$$\begin{aligned}D_{\text{KL}}(P||Q) &= \mathbb{E}_{X \sim P} \left[ \log \frac{P(X)}{Q(X)} \right] = \mathbb{E}_{X \sim P} [\log P(X) - \log Q(X)] \\ H(P, Q) &= -\mathbb{E}_{X \sim P} [\log Q(X)]\end{aligned}$$

---

<sup>1</sup>In fact, mean squared error is the KL divergence between the empirical distribution of the data and a Gaussian model.

These quantities are closely related - minimizing the cross-entropy  $H(P, Q)$  with respect to  $Q$  is equivalent to minimize the KL-divergence  $D_{\text{KL}}(P||Q)$ . Their relationship is given in the following equations

$$\begin{aligned} D_{\text{KL}}(P||Q) &= \mathbb{E}_{X \sim P}[\log P(X)] - \mathbb{E}_{X \sim P}[\log Q(X)] \\ &= H(P, Q) + \mathbb{E}_{X \sim P}[\log P(X)] \\ &= H(P, Q) - H(P) \\ \implies H(P, Q) &= H(P) + D_{\text{KL}}(P||Q) \end{aligned}$$

where Shannon entropy of a distribution  $P$  is defined as in class,

$$H(P) = -\mathbb{E}_{X \sim P}[\log P(X)]$$

## Question 2.

We are given the following information about the 2-layer neural network for  $K$ -class classification problem.

$$\begin{aligned} z^{[i]} &= W^{[i]} a^{[i-1]} + b^{[i]}, \quad i = 1, 2 \quad a^{[0]} = x \\ a^{[1]} &= \text{LEAKYRELU} \left( z^{[1]}, \alpha = 10^{-2} \right) \\ \hat{y} &= \text{SOFTMAX} \left( z^{[2]} \right) = \sigma \left( z^{[2]} \right) \quad (\text{say}) \\ \mathcal{L}(\hat{y}) &= - \sum_{i=1}^K y_i \log(\hat{y}_i) \end{aligned}$$

- (a) Let us assume  $z^{[2]}$  is of size  $K \times 1$  (since it is a  $K$ -class classification problem). Since  $z^{[1]}$  is of size  $D_a \times 1$ ,  $a^{[1]}$  must be of the same size, and hence, the  $W^{[2]}$  must be of size  $K \times D_a$ . Following this,  $b^{[2]}$  must be of size  $K \times 1$ .

$$\underbrace{z^{[2]}}_{K \times 1} = \underbrace{W^{[2]}}_{K \times D_a} \underbrace{a^{[1]}}_{D_a \times 1} + \underbrace{b^{[2]}}_{K \times 1}$$

By the same logic,  $W^{[1]}$  must be of size  $D_a \times D_x$ , since the output  $z^{[1]}$  must be of dimension  $D_a \times 1$  while the input is of size  $D_x \times 1$ .

Now, if we want to vectorize our input to a batch size of  $m$ , the input  $X$  will be of size  $D_x \times m$ . This means that the output  $z^{[1]}$  of the hidden layer will be  $D_a \times m$ , since

$$\underbrace{z^{[1]}}_{D_a \times m} = \underbrace{W^{[1]}}_{D_a \times D_x} \underbrace{X}_{D_x \times m} + \underbrace{b^{[1]}}_{D_a \times 1^*}$$

**\*Note:** Bias is added component-wise to each column of  $W^{[1]}X$  in case of vectorization of operations.

- (b) Let us denote the sum of exponentials (normalizing term) of the entries in  $z^{[2]}$  by

$$\mathbf{z} = \sum_{j=1}^K \exp z_j^{[2]}, \text{ which means } \hat{y} = \sigma \left( z^{[2]} \right) = \frac{\exp z^{[2]}}{\mathbf{z}}$$

We are required to find the partial derivative of the  $k^{\text{th}}$  entry in  $\hat{y}$  with respect to the  $k^{\text{th}}$  entry in  $z^{[2]}$ .

$$\begin{aligned} \frac{\partial \hat{y}_k}{\partial z_k^{[2]}} &= \frac{\partial}{\partial z_k^{[2]}} \left( \sigma \left( z_k^{[2]} \right) \right) = \frac{\partial}{\partial z_k^{[2]}} \left( \frac{\exp z_k^{[2]}}{\mathbf{z}} \right) \\ &= \frac{1}{\mathbf{z}^2} \left( \mathbf{z} \frac{\partial}{\partial z_k^{[2]}} \left( \exp z_k^{[2]} \right) - \exp z_k^{[2]} \frac{\partial}{\partial z_k^{[2]}} (\mathbf{z}) \right) \\ &= \frac{1}{\mathbf{z}^2} \left( \mathbf{z} \exp z_k^{[2]} - \left( \exp z_k^{[2]} \right)^2 \right) = \frac{\exp z_k^{[2]}}{\mathbf{z}} - \left( \frac{\exp z_k^{[2]}}{\mathbf{z}} \right)^2 \\ &= \hat{y}_k - \hat{y}_k^2 = \hat{y}_k (1 - \hat{y}_k) \end{aligned}$$

- (c) Next, we find the partial derivative of the  $k^{th}$  entry in  $\hat{y}$  with respect to the  $i^{th}$  entry in  $z^{[2]}$ , given  $i \neq k$ .

$$\begin{aligned}\frac{\partial \hat{y}_k}{\partial z_i^{[2]}} &= \frac{\partial}{\partial z_i^{[2]}} \left( \sigma \left( z_k^{[2]} \right) \right) = \frac{\partial}{\partial z_i^{[2]}} \left( \frac{\exp z_k^{[2]}}{\mathbf{z}} \right) \\ &= \frac{1}{\mathbf{z}^2} \left( \mathbf{z} \frac{\partial}{\partial z_i^{[2]}} \left( \exp z_k^{[2]} \right) - \exp z_k^{[2]} \frac{\partial}{\partial z_i^{[2]}} (\mathbf{z}) \right) \\ &= \frac{1}{\mathbf{z}^2} \left( 0 - \exp z_k^{[2]} \exp z_i^{[2]} \right) = -\frac{\exp z_k^{[2]}}{\mathbf{z}} \cdot \frac{\exp z_i^{[2]}}{\mathbf{z}} \\ &= -\hat{y}_k \hat{y}_i\end{aligned}$$

- (d) We are interested in the partial derivative of the loss with respect to entries in  $z^{[2]}$ . Using the chain rule, we have

$$\frac{\partial \mathcal{L}}{\partial z_i^{[2]}} = \frac{\partial \mathcal{L}}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial z_i^{[2]}}$$

With the above (given) loss function, we can find

$$\frac{\partial \mathcal{L}}{\partial \hat{y}_i} = \frac{\partial}{\partial \hat{y}_i} \left( -\sum_{j=1}^K y_j \log(\hat{y}_j) \right) = -\sum_{j=1}^K y_j \frac{\partial}{\partial \hat{y}_i} (\log(\hat{y}_j)) = -\frac{y_i}{\hat{y}_i}$$

Note that as a result from **Question 2.** (b),

$$\frac{\partial \hat{y}_i}{\partial z_i^{[2]}} = \hat{y}_i(1 - \hat{y}_i)$$

So, we have

$$\frac{\partial \mathcal{L}}{\partial z_i^{[2]}} = -\frac{y_i}{\hat{y}_i} \cdot \hat{y}_i(1 - \hat{y}_i) = -y_i(1 - \hat{y}_i)$$

Now, we are given that  $y_k = 1$  and  $y_i = 0, i \neq k$ . So, we get (where  $\mathbb{I}$  is an indicator variable)

$$\frac{\partial \mathcal{L}}{\partial z_i^{[2]}} = \mathbb{I}_{\{i=k\}}(\hat{y}_i - 1) = \begin{cases} \hat{y}_k - 1 & \text{if } i = k \\ 0 & \text{if } i \neq k \end{cases}$$

- (e) While implementing the softmax function, we run into the problem of numerical overflow and underflow. Overflow occurs while exponentiating large numbers, when the result exceeds the computer's largest representable float value. In such cases, the result of  $e^z$  is estimated to be infinity (`torch.inf` in PyTorch). Conversely, underflow occurs while exponentiating small numbers, when extremely small values are represented by zeros. Rarely, it is also possible that the denominator is extremely small, which makes division unstable.

A common modification to resolve the issue is normalizing the input by max-shifting. In this technique, we subtract the maximum of the vector from all its entries, making all of them non-positive, and at least one entry 0. The output probabilities do not change because the relative difference between the entries in  $z$  does not change, but the numerical operations become more stable. So, for a vector  $z$  of size  $m$ , we get

$$\begin{aligned}z^* &= z - \max_{k=1,2,\dots,m} z_k \\ \sigma(z) &= \sigma(z^*) = \frac{\exp z^*}{\sum_{j=1}^m \exp z_j^*}\end{aligned}$$

## Image Classification

The solutions to the questions in this section are given in `classification.ipynb`.

### Question 1.

- (a) The dataset was downloaded to my system from the Google Drive (I know, i couldn't figure out how to connect it to Google Colab and use GPU :sad:). The dataset was loaded using a custom class, `WildlifeDataset`, which also handled the category to label mapping.

- (b) The splits of ratio 0.7:0.1:0.2 were generated and loaded using dataloaders with batch size 256.
- (c) To visualize the data distribution across class labels, we count the number of images in the entire dataset. Then, I compared the proportion of images in each class in the training and validation set. The plots are given in the following figures.

## Question 2.

- (a) The given CNN architecture was constructed in the class `ConvNet`. The model constructed using PyTorch.
- (b) I used wandb to log the training and validation losses and accuracies per epoch. The plots are given in the figures below.
- (c) Looking at the figures, it is clear that the model is overfitting after training for a few epochs, since the validation loss stops decreasing sharply while the training loss keeps decreasing. The final training accuracy after 15 epochs is close to 0.86+, while the validation accuracy was around 0.58+ only.
- (d) The confusion matrix, logged using wandb, is given in the following figure. However, it is still surprising that the accuracy on the test set was 0.91+. This is a good example of why accuracy is not the only indicator of a model's performance.

METRIC	VALUE
Accuracy	0.91773+
Precision	0.58870+
Recall	0.58870+
F1-Score	0.58870+

Table 1: Evaluation Metrics for the given CNN architecture

- (e) Finally, we visualize 3 misclassified images from each class. These are also given in the same notebook. A few examples are also given in the following figure. There were quite a few examples where the ground truth class was not present or did not cover a majority of the image. For example, some images of class `amur_leopard` were not clear. In some cases, the ground truth class was also not present, for example, in case of `black_bear`. There was also a case where the image looks like the predicted class - `brown_bear` being predicted as `black_bear` and `wild_bear`.

## Question 3.

- (a) Next, we finetune the Resnet-18 model available in PyTorch on the Russian Wildlife dataset and again log the accuracies and losses using wandb.
- (b) This model is not overfitting as much as the previous model. This is inferred because the difference between the training loss and validation loss is lower, and the same goes for the accuracy.
- (c) The evaluation metrics for the Resnet-18 architecture are given in the table below.

METRIC	VALUE
Accuracy	0.96937+
Precision	0.84688+
Recall	0.84688+
F1-Score	0.84688+

Table 2: Evaluation Metrics for the Resnet-18 architecture

- (d) The 2D and 3D representations of the feature space generated by input samples from training and validation sets by the classifier were generated and are given in the following figures.

## Question 4.

- (a) To perform data augmentation to increase robustness of the model, I used different transforms to add randomness in the data, which are available in `torchvision.transforms`. Specifically, I applied random affine transforms, random horizontal flipping with probability 0.5 for each image, and added random color noise or jitter in the images, altering their color schemes (brightness, contrast, saturation, and hue). A total of 2000 augmented images were added to the dataset.

- (b) As asked in the Question, the same steps as **Question 2.3 (a)** were followed to train the Resnet-18 model on the augmented dataset.
- (c) Looking at the loss curves, it is clear that this model performs even better than before, and is not overfitting. The curves are given in the following figures.
- (d) The evaluation metrics for the model trained on the augmented dataset are given in the following table.

METRIC	VALUE
Accuracy	0.97653+
Precision	0.88265+
Recall	0.0.88265+
F1-Score	0.88265+

Table 3: Evaluation Metrics for the Resnet-18 architecture on augmented dataset

### Question 5.

The analysis for this question is also given in the notebook. The euclidean distances between the feature space representations of the misclassified images were calculated from the mean of both the ground truth and predicted classes. It is interesting to note that in some cases, the distance to the mean of the ground truth class was lower than the distance to the mean of the predicted class. However, such cases were rare. A few examples are given in the following figures.

### Question 6.

On comparing all the three models, it is pretty clear that the vanilla CNN architecture performs well, but not as well as we would like. Its recall and F1-scores are low. This is due to its comparatively lower capacity and higher tendency to overfit the data. On the flip side, the Resnet-18 architectures performed extremely well. Both models did not overfit, achieving high scores for all common evaluation metrics. The model trained with the augmented dataset outperformed the other one by a small margin, since adding some noise to the dataset increases its ability to generalize.

## Image Segmentation

The solutions to the questions in this section are given in two notebooks, namely **segmentation-1.ipynb** and **segmentation-2.ipynb** (By this time, I figured out connecting Google Drive to Google Colab, and hence the second notebook performed inferences using the GPU).

### Question 1.

- (a) Similar to the image classification problem, I created a custom dataset class, **IDD**, which also handles label mappings and mask color mappings for both the **IDD** and **Cityscapes** dataset. Since the images were larger in size, I used a batch size of 8 to create the dataloader. The dataset was imported from Google Drive :)
- (b) To visualize the data distribution across the dataset, I used pixel counts in the mask - the frequency of pixels belonging to each of the classes. Due to computational limitations, the plot was generated only for a fraction of the dataset. The distribution is given in the following figure.
- (c) For each class, I visualize the images and their masks (and show which class is being highlighted in each case). The masks were color coded according to the conventions of the **IDD** Dataset. The results are given in **segmentation-1.ipynb**, and a few examples are given in the figure below.

### Question 2.

- (a) The **DeepLabv3Plus-PyTorch** Resnet-101 model was used for performing inferences on 30% **IDD** dataset (using it as test set). The inference was performed in the second notebook, which is GPU enabled.
- (b) The classwise performances in terms of the given metrics were evaluated. The results are given in the following table (Forgive me, I did not have enough time to present these in a better manner).

- (c) I visualize three images per class where the IoU between the predicted and true masks is at most 0.5. It was easy to notice the faults - in some cases, the IoU was close to 0.48+, for example, in case of **Road**. In other cases, like **Sidewalk** or **Traffic-Light**, the predicted masks did not contain these classes (hence the IoUs were 0). This is probably because the model was trained on a different and less complex dataset, resulting in classes that occur on only a few pixels per image to be *ignored*. In such cases, the surrounding objects and environment were overpowering the ground truth class. For example, most pixels of **Sidewalk** were being predicted as **Road**.

### Question 3.

- (a) The confusion matrix for the final segmentation was generated and is given in the following figure. We can infer that classes like **Road** and **Sky** are being predicted well, but others were not - this is seen due to the low count of true positives on the diagonals corresponding to these classes. It is also noticeable that most mistakes are being made in less frequent classes, which are being predicted as classes which are more prevalent.
- (b) The classwise precision, recall, and F1-scores were calculated and are given in the following table (again, owing to the lack of time to present the results better). Again, it is clear that the model needs improvements in classes which occur on a small fraction of an image, for example classes like

### Question 4.

- (a) The GPU enabled notebook was used to perform inference on the Cityscapes dataset. Like before, the generated masks were stored in Google Drive.
- (b) The following figure shows the confusion matrix for these inferences, which are arguably better than the previous case. This is because the model was trained on Cityscapes dataset, and is hence able to recognize classes better. Now, more classes have a higher count of true positives (diagonal entries in the confusion matrix).
- (c) Similar to before, the worst performing class performs bad because it is being dominated by their environments containing classes like **Road**, **Vegetation**, and **Sky** - which span a larger area. Exactly these dominating classes perform well, since they were being predicted *more often* in a sense, and covering most of their true pixels.